

```

#!pip install numpy==1.26.4
#!pip install scikit-learn-extra

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn_extra.cluster import KMedoids
from sklearn.cluster import KMeans

print(np.__version__)

1.26.4

df = pd.read_csv('train.csv')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 878049 entries, 0 to 878048
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   Dates            878049 non-null object
1   Category         878049 non-null object
2   Descript         878049 non-null object
3   DayOfWeek        878049 non-null object
4   PdDistrict       878049 non-null object
5   Resolution       878049 non-null object
6   Address          878049 non-null object
7   X                878049 non-null float64
8   Y                878049 non-null float64
dtypes: float64(2), object(7)
memory usage: 60.3+ MB

df.isna().sum()

Dates            0
Category         0
Descript         0
DayOfWeek        0
PdDistrict       0
Resolution       0
Address          0
X                0
Y                0
dtype: int64

df.dropna(inplace=True)

df.duplicated().sum()

```

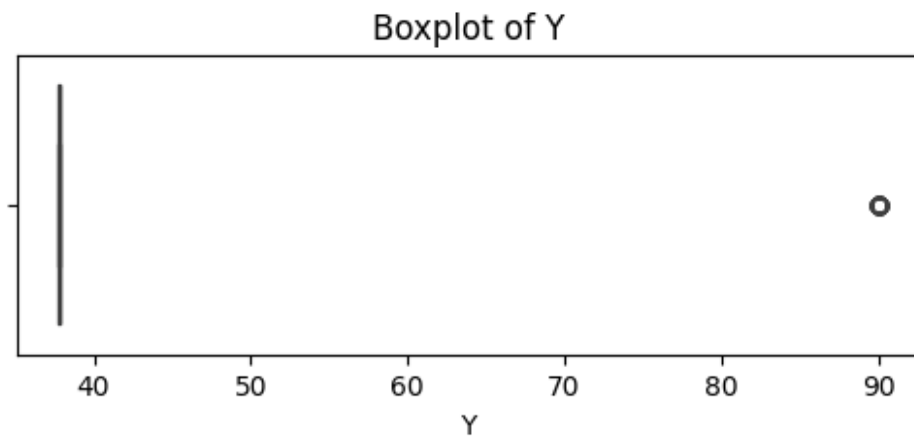
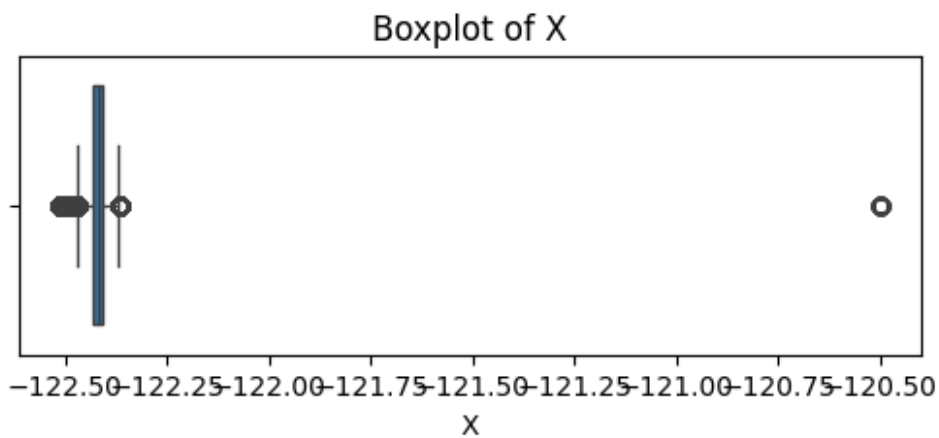
2323

```
df.drop_duplicates(inplace = True)


df['Dates'] = pd.to_datetime(df['Dates'])

df['year'] = df['Dates'].dt.year
df['month'] = df['Dates'].dt.month
df['day'] = df['Dates'].dt.day
df['hour'] = df['Dates'].dt.hour


numerical_cols = ['X', 'Y', 'hour', 'day', 'month']
for col in numerical_cols:
    plt.figure(figsize=(6, 2))
    sns.boxplot(x=df[col])
    plt.title(f"Boxplot of {col}")
    plt.show()
```




Boxplot of hour



Boxplot of day



Boxplot of month



```
df = df[(df['Y'] < 38) & (df['Y'] > 37) & (df['X'] < -122) & (df['X'] > -123)]
```

```
plt.figure(figsize=(10, 5), constrained_layout=True)
```

```

# Get top 10 categories
tmp = df['Category'].value_counts().head(10)

# Create a colored barplot
sns.barplot(x=tmp.index, y=tmp.values, palette='husl') # You can
change 'viridis' to any other palette

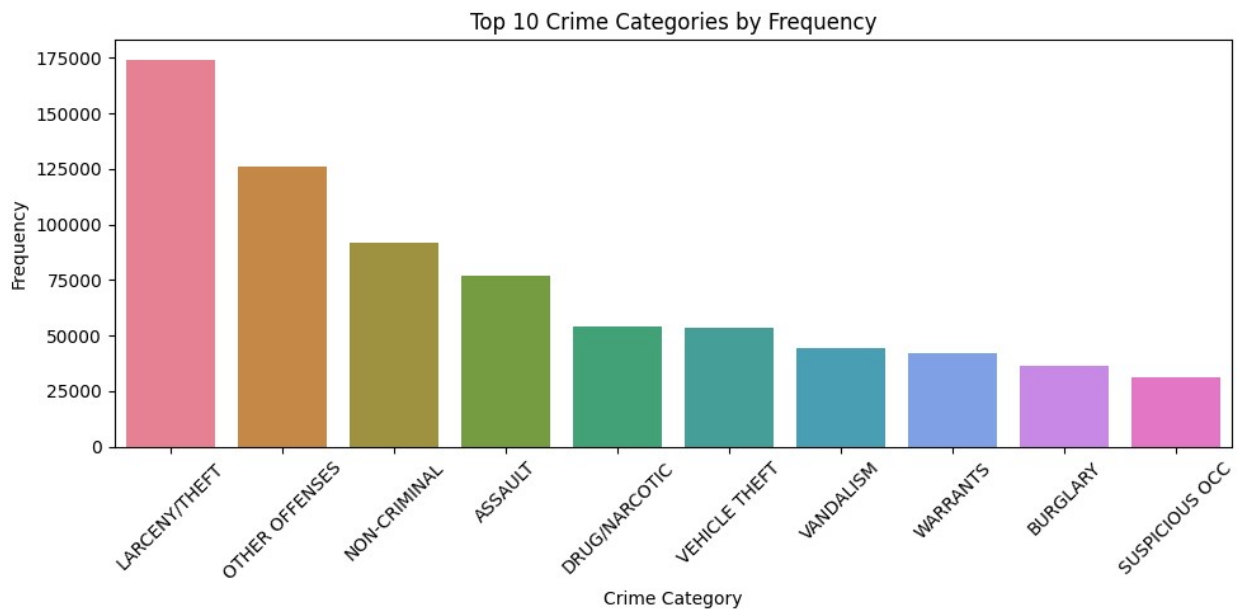
plt.title('Top 10 Crime Categories by Frequency')
plt.xticks(rotation=45)
plt.xlabel('Crime Category')
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()

<ipython-input-67-2225c78d2a53>:7: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be
removed in v0.14.0. Assign the `x` variable to `hue` and set
`legend=False` for the same effect.

sns.barplot(x=tmp.index, y=tmp.values, palette='husl') # You can
change 'viridis' to any other palette
<ipython-input-67-2225c78d2a53>:13: UserWarning: The figure layout has
changed to tight
plt.tight_layout()

```



```

dates_indexed = df.set_index('Dates')

category_trend = dates_indexed.groupby(
    [pd.Grouper(freq='M'), 'Category']).size().unstack(fill_value=0)

```

```

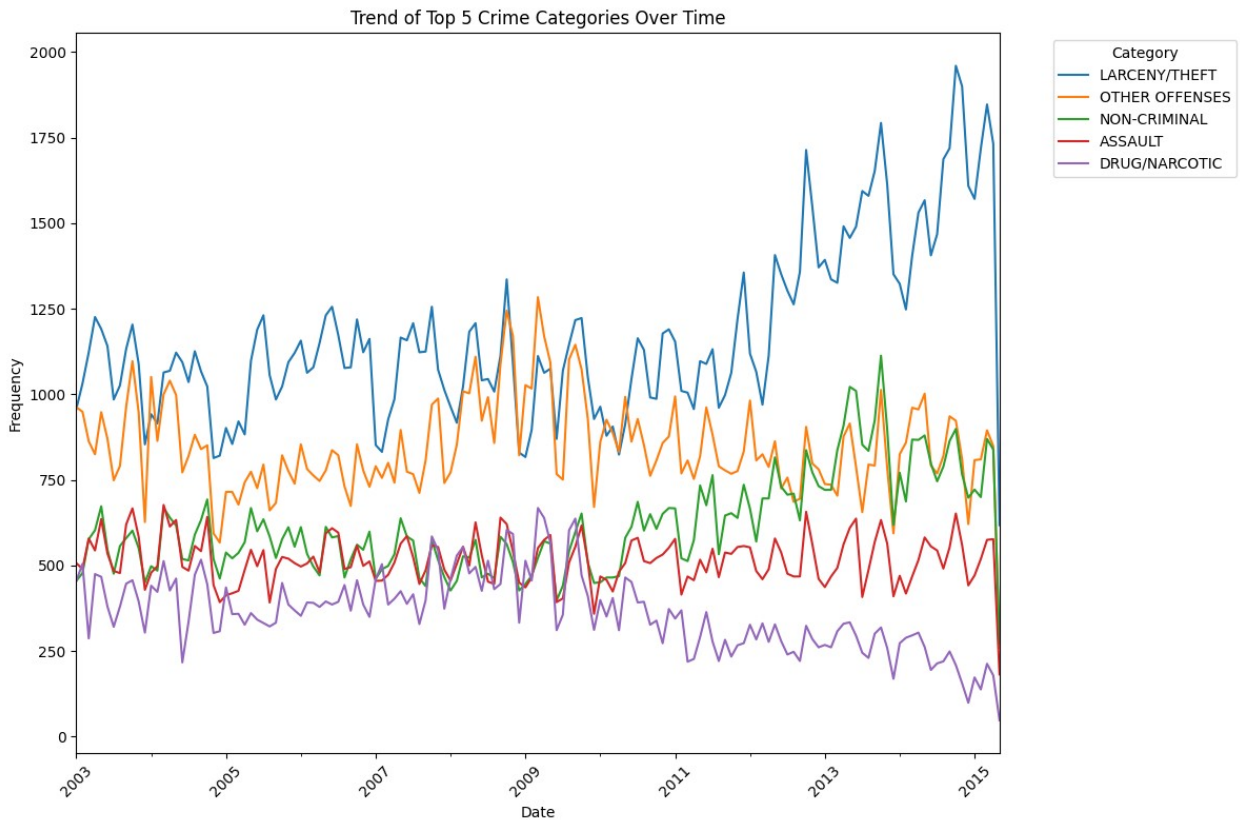
# Get the top N categories by total count
top_categories =
dates_indexed['Category'].value_counts().nlargest(5).index

# Filter the category trend data to only include top categories
top_category_trend = category_trend[top_categories]

# Plot only the top categories
top_category_trend.plot(kind='line', figsize=(12, 8))
plt.title('Trend of Top 5 Crime Categories Over Time')
plt.ylabel('Frequency')
plt.xlabel('Date')
plt.xticks(rotation=45)
plt.legend(title='Category', bbox_to_anchor=(1.05, 1), loc='upper
left')
plt.tight_layout()
plt.show()

<ipython-input-68-784c25f3e55a>:4: FutureWarning: 'M' is deprecated
and will be removed in a future version, please use 'ME' instead.
[pd.Grouper(freq='M'), 'Category']).size().unstack(fill_value=0)

```



```

sns.barplot(
    x=df.DayOfWeek.value_counts().index,

```

```

y=df.DayOfWeek.value_counts().values,
order=['Sunday', 'Monday', 'Tuesday', 'Wednesday', 'Thursday',
'Friday', 'Saturday'],
palette='husl'
)


```

<ipython-input-69-4d25189fd6cd>:1: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(
```

<Axes: xlabel='DayOfWeek'>



```

# Group by hour and category, count occurrences
hourly_category_trend = df.groupby(['hour',
'Category']).size().unstack(fill_value=0)

# Plot the category trends by hour
plt.figure(figsize=(12, 8))
hourly_category_trend.plot(kind='line', figsize=(14, 8))
plt.title('Crime Frequency by Category and Hour of the Day')


```

```

plt.ylabel('Frequency')
plt.xlabel('Hour of the Day')
plt.xticks(ticks=range(0, 24), labels=range(0, 24)) # Ensure x-axis
has all hours
plt.legend(title='Category', bbox_to_anchor=(1.05, 1), loc='upper
left')
plt.tight_layout()
plt.show()

```

<Figure size 1200x800 with 0 Axes>



```

district_crime_counts = df.groupby(['PdDistrict',
'Category']).size().reset_index(name='Count')

# Sort and pick top 3 crimes per district
top3_crimes_per_district = (
    district_crime_counts
    .sort_values(['PdDistrict', 'Count'], ascending=[True, False])
    .groupby('PdDistrict')
    .head(3)
)

# Now plot
plt.figure(figsize=(15, 10))
sns.barplot(data=top3_crimes_per_district, x='PdDistrict', y='Count',


```

```

hue='Category')

plt.title('Top 3 Crimes in Each Police District')
plt.xlabel('Police District')
plt.ylabel('Number of Crimes')
plt.xticks(rotation=45)
plt.legend(title='Crime Category', bbox_to_anchor=(1.05, 1),
loc='upper left')
plt.tight_layout()
plt.show()

```



```

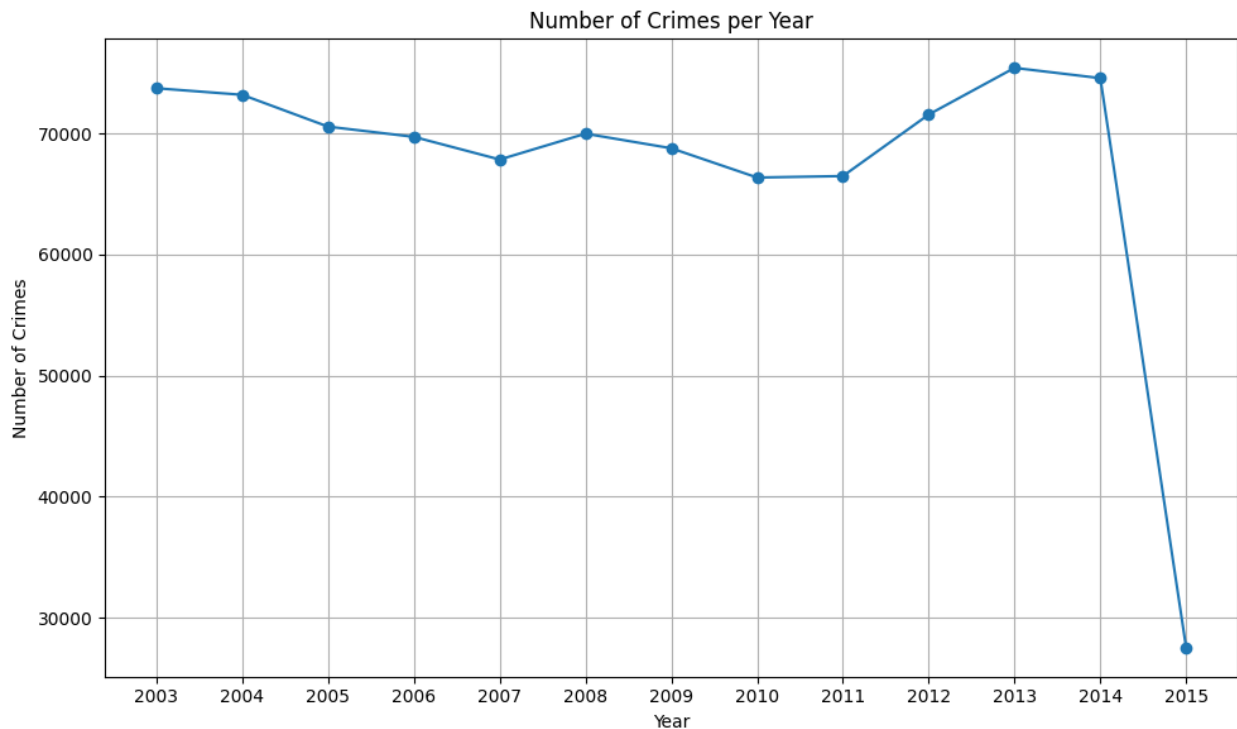
crime_counts_by_year = df.groupby('year').size()

# Plot
plt.figure(figsize=(10, 6))
plt.plot(crime_counts_by_year.index, crime_counts_by_year.values,
marker='o')
plt.title('Number of Crimes per Year')
plt.xlabel('Year')
plt.ylabel('Number of Crimes')
plt.grid(True)
plt.xticks(crime_counts_by_year.index) # show all years on x-axis

```



```
plt.tight_layout()
plt.show()
```

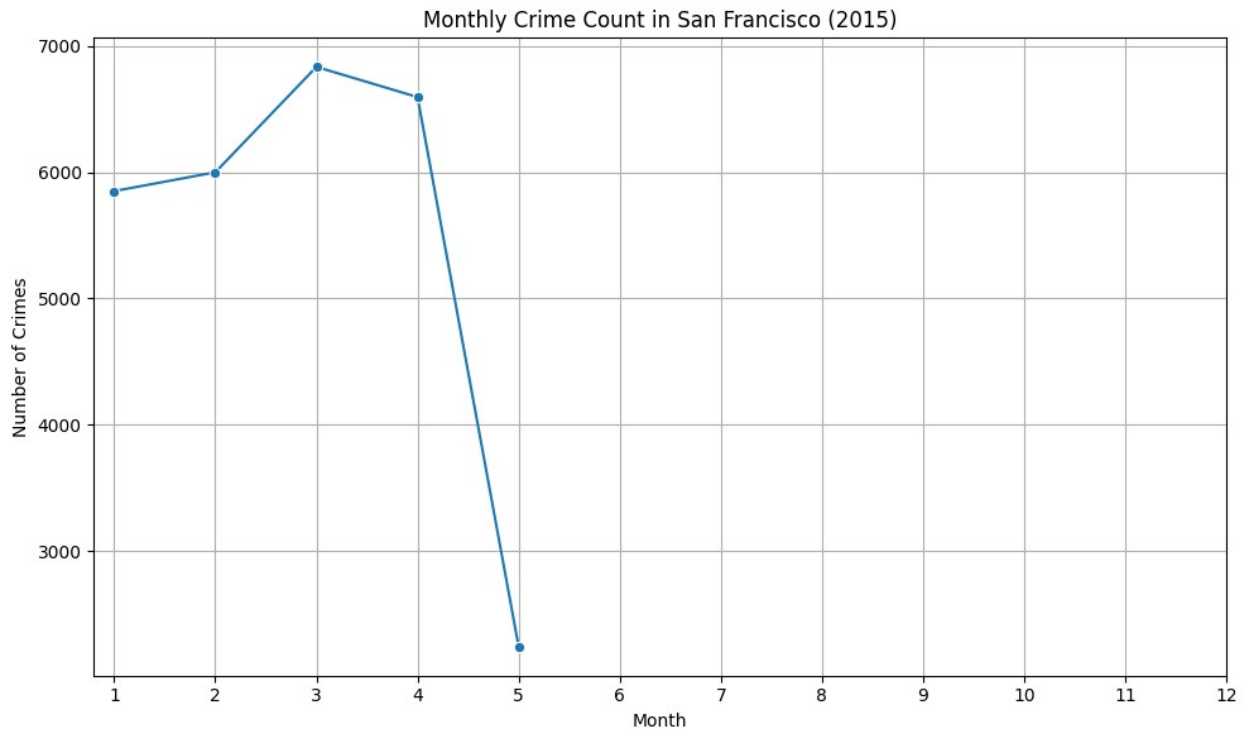


```
# Filter for 2015
df_2015 = df[df['Dates'].dt.year == 2015]

# Group by month
monthly_crimes_2015 =
df_2015.groupby(df_2015['Dates'].dt.month).size().reset_index(name='crime_count')

# Plot
plt.figure(figsize=(10, 6))
sns.lineplot(data=monthly_crimes_2015, x='Dates', y='crime_count',
marker='o')

plt.title('Monthly Crime Count in San Francisco (2015)')
plt.xlabel('Month')
plt.ylabel('Number of Crimes')
plt.xticks(range(1, 13))
plt.grid(True)
plt.tight_layout()
plt.show()
```



```
from sklearn.preprocessing import MinMaxScaler
from sklearn.preprocessing import LabelEncoder
```

```
df1=df[(df["year"]==2015)]
```

```
sc=MinMaxScaler()
sc.fit(df1[["X"]])
df1["X_scale"]=sc.transform(df1[["X"]])
```

```
sc.fit(df1[["Y"]])
df1["Y_scale"]=sc.transform(df1[["Y"]])
```

```
le=LabelEncoder()
df1["Category"]=le.fit_transform(df1["Category"])
```

```
sc.fit(df1[["Category"]])
df1["cat_sc"]=sc.transform(df1[["Category"]])
df1.head()
```

```
<ipython-input-74-db92acb92d79>:10: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df1["X_scale"]=sc.transform(df1[["X"]])
<ipython-input-74-db92acb92d79>:14: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df1["Y_scale"]=sc.transform(df1[["Y"]])
<ipython-input-74-db92acb92d79>:17: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df1["Category"]=le.fit_transform(df1["Category"])
<ipython-input-74-db92acb92d79>:20: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df1["cat_sc"]=sc.transform(df1[["Category"]])

{"summary": "{\n  \"name\": \"df1\",\n  \"rows\": 27522,\n  \"fields\": [\n    {\n      \"column\": \"Dates\",\n      \"properties\": {\n        \"dtype\": \"date\",\n        \"min\": \"2015-01-05 00:01:00\",\n        \"max\": \"2015-05-13 23:53:00\",\n        \"num_unique_values\": 12096,\n        \"samples\": [\n          \"2015-05-11 11:10:00\",\n          \"2015-01-07 14:02:00\",\n          \"2015-05-03 23:03:00\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Category\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 9,\n        \"min\": 0,\n        \"max\": 36,\n        \"num_unique_values\": 37,\n        \"samples\": [\n          19,\n          7,\n          33\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Descript\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 559,\n        \"samples\": [\n          \"VANDALISM OR GRÄFFITI TOOLS, POSSESSION\",\n          \"INTOXICATED JUVENILE\",\n          \"BURGLARY,BLDG. UNDER CONSTRUCTION, UNLAWFUL ENTRY\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"DayOfWeek\",\n      \"properties\": {\n        \"dtype\": \"category\",\n        \"num_unique_values\": 7,\n        \"samples\": [\n
```

```

\ "Wednesday\ ",\n          \ "Tuesday\ ",\n          \ "Friday\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "PdDistrict\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "category\ ",\n          \ "num_unique_values\ ": 10,\n          \ "samples\ ": [\n          \ "MISSION\ ",\n          \ "PARK\ ",\n          \ "CENTRAL\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "Resolution\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "category\ ",\n          \ "num_unique_values\ ": 10,\n          \ "samples\ ": [\n          \ "CLEARED-CONTACT JUVENILE FOR MORE INFO\ ",\n          \ "NONE\ ",\n          \ "UNFOUNDED\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "Address\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "category\ ",\n          \ "num_unique_values\ ": 8115,\n          \ "samples\ ": [\n          \ "900 Block of SCOTT ST\ ",\n          \ "BEACH ST / STOCKTON ST\ ",\n          \ "14TH ST / MINNA ST\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "X\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "number\ ",\n          \ "std\ ": 0.026236931909153344,\n          \ "min\ ": -122.513642064265,\n          \ "max\ ": -122.365565425353,\n          \ "num_unique_values\ ": 8885,\n          \ "samples\ ": [\n          \ "122.470873483267\ ",\n          \ "122.478575817534\ ",\n          \ "122.389487749378\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "Y\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "number\ ",\n          \ "std\ ": 0.024450013487924846,\n          \ "min\ ": 37.7080829769335,\n          \ "max\ ": 37.819923463743,\n          \ "num_unique_values\ ": 8885,\n          \ "samples\ ": [\n          \ "37.7492742889732\ ",\n          \ "37.7850879697075\ ",\n          \ "37.724599376829296\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "year\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "int32\ ",\n          \ "num_unique_values\ ": 1,\n          \ "samples\ ": [\n          \ "2015\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "month\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "int32\ ",\n          \ "num_unique_values\ ": 5,\n          \ "samples\ ": [\n          \ "4\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "day\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "int32\ ",\n          \ "num_unique_values\ ": 30,\n          \ "samples\ ": [\n          \ "23\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "hour\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "int32\ ",\n          \ "num_unique_values\ ": 24,\n          \ "samples\ ": [\n          \ "15\ "\n          ],\n          \ "semantic_type\ ": \ "\",\n          \ "description\ ": \ "\",\n          \ "column\ ": \ "X_scale\ ",\n          \ "properties\ ": {\n          \ "dtype\ ": \ "number\ ",\n          \ "std\ ": 0.1771848152546543,\n          \ "min\ ": 0.0,\n          \ "max\ ": 1.0,\n          \ "num_unique_values\ ": 8885,\n

```

```

{"samples\":[\n          0.2888273350356485\n        ],\n "semantic_type\":"\", \n "description\":"\", \n\n    },\n    {\n      \"column\":"Y_scale",\n      \"properties\":\n        {\n          \"dtype\":"number",\n          \"std\":\n            0.21861504885587024,\n          \"min\": 0.0,\n          \"max\": 1.0,\n          \"num_unique_values\": 8885,\n          \"samples\":[\n            0.368304119686627\n          ],\n          \"semantic_type\":"\", \n          \"description\":"\", \n        },\n        {\n          \"column\":"cat_sc",\n          \"properties\":\n            {\n              \"dtype\":"number",\n              \"std\": 0.2622990554348488,\n              \"min\": 0.0,\n              \"max\": 1.0,\n              \"num_unique_values\": 37,\n              \"samples\":[\n                0.5277777777777778\n              ],\n              \"semantic_type\":"\", \n              \"description\":"\", \n            }\n          }\n        }\n      ],\n      \"type\":\"dataframe\", \"variable_name\":\"df1\"}

```

```

from scipy.cluster.hierarchy import linkage, dendrogram, fcluster
import matplotlib.pyplot as plt
import pandas as pd

```

```

# 1. Take a random sample (e.g., 500 rows)
sample_df = df1.sample(n=250, random_state=42) # Set random_state for reproducibility

```

```

# 2. Prepare the data for clustering
X = sample_df[['X_scale', 'Y_scale', 'cat_sc']]

```

```

# 3. Compute the linkage matrix
linked = linkage(X, method='ward')

```

```

# 4. Plot the dendrogram
num_clusters = 3
plt.figure(figsize=(16, 8))
dendrogram(linked,
            orientation='top',
            distance_sort='descending',
            show_leaf_counts=True,
            color_threshold=linked[-num_clusters, 2],
            no_labels=False)
plt.axhline(y=linked[-num_clusters, 2], color='red', linestyle='--')
plt.title("Hierarchical Clustering Dendrogram (Sample of 500)")
plt.xlabel("Sample Index")
plt.ylabel("Distance")
plt.tight_layout()
plt.show()

```

```

# 5. Assign cluster labels
sample_df['hier_cluster'] = fcluster(linked, num_clusters,
criterion='maxclust')

```


```

# Optional: Add a distance-based superclass

```

```
sample_df['superclass'] = fcluster(linked, t=12, criterion='distance')

# 6. Display result
print(sample_df[['X_scale', 'Y_scale', 'cat_sc', 'hier_cluster',
'superclass']].head())
```



	X_scale	Y_scale	cat_sc	hier_cluster	superclass
24360	0.744461	0.602087	0.444444	1	1
18715	0.684729	0.665889	0.027778	1	1
13348	0.715824	0.616021	0.444444	1	1
22221	0.834068	0.727802	0.444444	1	1
10046	0.402816	0.075044	0.027778	3	1

```
from sklearn.metrics import silhouette_score
```

```
# Use the same feature matrix X
```

```
score = silhouette_score(X, sample_df['hier_cluster'])
```

```
print(f"Silhouette Score for Hierarchical Clustering  
(k={num_clusters}): {score:.4f}")
```

```
Silhouette Score for Hierarchical Clustering (k=3): 0.3069
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 875659 entries, 0 to 878048
```

```
Data columns (total 13 columns):
```


#	Column	Non-Null Count	Dtype
0	Dates	875659 non-null	datetime64[ns]
1	Category	875659 non-null	object

```
2   Descript      875659 non-null  object
3   DayOfWeek     875659 non-null  object
4   PdDistrict    875659 non-null  object
5   Resolution    875659 non-null  object
6   Address       875659 non-null  object
7   X             875659 non-null  float64
8   Y             875659 non-null  float64
9   year          875659 non-null  int32
10  month         875659 non-null  int32
11  day           875659 non-null  int32
12  hour          875659 non-null  int32
dtypes: datetime64[ns](1), float64(2), int32(4), object(6)
memory usage: 80.2+ MB
```

```
from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

best = []
for i in range(1, 15):
    km1 = KMeans(n_clusters=i, random_state=42)
    km1.fit(df[["X", "Y"]]) # Combine X and Y features
    best.append(km1.inertia_)

plt.plot(range(1, 15), best, marker='o')
plt.xlabel("Number of Clusters (K)")
plt.ylabel("Inertia (Within-cluster Sum of Squares)")
plt.title("Elbow Method for Optimal K")
plt.show()
```



```
!pip install scikit-learn-extra
```

```
Requirement already satisfied: scikit-learn-extra in  
/usr/local/lib/python3.11/dist-packages (0.3.0)  
Requirement already satisfied: numpy>=1.13.3 in  
/usr/local/lib/python3.11/dist-packages (from scikit-learn-extra)  
(1.26.4)  
Requirement already satisfied: scipy>=0.19.1 in  
/usr/local/lib/python3.11/dist-packages (from scikit-learn-extra)  
(1.15.2)  
Requirement already satisfied: scikit-learn>=0.23.0 in  
/usr/local/lib/python3.11/dist-packages (from scikit-learn-extra)  
(1.6.1)  
Requirement already satisfied: joblib>=1.2.0 in  
/usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.23.0-  
>scikit-learn-extra) (1.4.2)  
Requirement already satisfied: threadpoolctl>=3.1.0 in  
/usr/local/lib/python3.11/dist-packages (from scikit-learn>=0.23.0-  
>scikit-learn-extra) (3.6.0)
```

```
#Kmedoids when K = 5
```



```

k=5

# Combine features into a single matrix
X = df1[['X', 'Y']]

# Fit KMedoids with the combined feature matrix
kmedoids = KMedoids(n_clusters=k).fit(X)

clusters = kmedoids.cluster_centers_
labels = kmedoids.labels_

/usr/local/lib/python3.11/dist-packages/sklearn_extra/cluster/_k_medoids.py:329: UserWarning: Cluster 2 is empty!
self.labels_[self.medoid_indices[2]] may not be labeled with its
corresponding cluster (2).
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn_extra/cluster/_k_medoids.py:329: UserWarning: Cluster 3 is empty!
self.labels_[self.medoid_indices[3]] may not be labeled with its
corresponding cluster (3).
  warnings.warn(
/usr/local/lib/python3.11/dist-packages/sklearn_extra/cluster/_k_medoids.py:329: UserWarning: Cluster 4 is empty!
self.labels_[self.medoid_indices[4]] may not be labeled with its
corresponding cluster (4).
  warnings.warn(

# Drop unwanted columns
df_2 = df1.drop(["Descript", "Resolution", "Address", "PdDistrict",
"DayOfWeek"], axis=1)

# Assign the cluster labels from K-Medoids
df_2['cluster'] = kmedoids.labels_

import plotly.express as px
figure1 = px.scatter_mapbox(df_2,
                           lat="Y", lon="X", # Latitude and
Longitude columns
                           color="cluster", # Color by cluster
                           hover_data=["Category", "X", "Y",
"cluster"], # Hover information
                           title="K-Medoids Clustering of Crime
Locations",
                           center=dict(lat=37.8, lon=-122.4), # Set
center for map (adjust if needed)
                           zoom=9, # Zoom level for the map
                           width=1100,
                           height=700,
                           mapbox_style="open-street-map") # Map

```

```
style
```

```
# Show the plot
```

```
figure1.show()
```

```
from sklearn.metrics import silhouette_score
```

```
sil_score = silhouette_score(X, labels)
```

```
print(f"Silhouette Score: {sil_score:.4f}")
```

```
Silhouette Score: 0.3984
```

```
from sklearn.metrics import davies_bouldin_score
```

```
db_score = davies_bouldin_score(X, labels)
```

```
print(f"Davies-Bouldin Index: {db_score:.4f}")
```

```
Davies-Bouldin Index: 0.8651
```

```
df_2.head()
```

```
{"summary":{"\n  \"name\": \"df_2\",\n  \"rows\": 27522,\n  \"fields\": [\n    {\n      \"column\": \"Dates\",\n      \"properties\": {\n        \"dtype\": \"date\",\n        \"min\": \"2015-01-05 00:01:00\",\n        \"max\": \"2015-05-13 23:53:00\",\n        \"num_unique_values\": 12096,\n        \"samples\": [\n          \"2015-05-11 11:10:00\",\n          \"2015-01-07 14:02:00\",\n          \"2015-05-03 23:03:00\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Category\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 9,\n        \"min\": 0,\n        \"max\": 36,\n        \"num_unique_values\": 37,\n        \"samples\": [\n          19,\n          7,\n          33\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"X\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.026236931909153344,\n        \"min\": -122.513642064265,\n        \"max\": -122.365565425353,\n        \"num_unique_values\": 8885,\n        \"samples\": [\n          122.470873483267,\n          -122.478575817534,\n          122.389487749378\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Y\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.024450013487924846,\n        \"min\": 37.7080829769335,\n        \"max\": 37.819923463743,\n        \"num_unique_values\": 8885,\n        \"samples\": [\n          37.7492742889732,\n          37.7850879697075,\n          37.724599376829296\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"year\",\n      \"properties\": {\n        \"dtype\": \"int32\",\n        \"num_unique_values\": 1,\n        \"samples\": [\n          2015\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    ]\n  }\n}
```

```

{"column": "month", "properties": {"dtype": "int32", "num_unique_values": 5, "samples": [4]}, "semantic_type": "", "description": ""}, {"column": "day", "properties": {"dtype": "int32", "num_unique_values": 30, "samples": [23]}, "semantic_type": "", "description": ""}, {"column": "hour", "properties": {"dtype": "int32", "num_unique_values": 24, "samples": [15]}, "semantic_type": "", "description": ""}, {"column": "X_scale", "properties": {"dtype": "number", "std": 0.1771848152546543, "min": 0.0, "max": 1.0, "num_unique_values": 8885, "samples": [0.2888273350356485]}, "semantic_type": "", "description": ""}, {"column": "Y_scale", "properties": {"dtype": "number", "std": 0.21861504885587024, "min": 0.0, "max": 1.0, "num_unique_values": 8885, "samples": [0.368304119686627]}, "semantic_type": "", "description": ""}, {"column": "cat_sc", "properties": {"dtype": "number", "std": 0.2622990554348488, "min": 0.0, "max": 1.0, "num_unique_values": 37, "samples": [0.5277777777777778]}, "semantic_type": "", "description": ""}, {"column": "cluster", "properties": {"dtype": "number", "std": 1, "min": 0, "max": 4, "num_unique_values": 5, "samples": [0]}, "semantic_type": "", "description": ""}]
{"type": "dataframe", "variable_name": "df_2"}

```

$$\text{Log Loss} = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K y_{i,k} \cdot \log(p_{i,k})$$

$$\hat{y}_i^{(1)} = \hat{y}_i^{(0)} + \eta \cdot f_1(x_i)$$

```

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from xgboost import XGBClassifier
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

def train_and_evaluate_xgboost(df):
    X = df.drop(columns=['cluster', 'Dates'])
    y = df['cluster']

    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

    model = XGBClassifier(eval_metric='mlogloss', random_state=42)
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    print(f" Accuracy: {acc * 100:.2f}%\n")

    print(" Classification Report:")
    print(classification_report(y_test, y_pred))

    plt.figure(figsize=(12, 8))
    sns.heatmap(confusion_matrix(y_test, y_pred), annot=False,
fmt='d', cmap='Blues')
    plt.title("Confusion Matrix", fontsize=16)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

    return model

model = train_and_evaluate_xgboost(df_2)


Accuracy: 99.76%

Classification Report:

```

	precision	recall	f1-score	support

0	1.00	1.00	1.00	2363
1	0.99	0.99	0.99	498
2	0.99	0.99	0.99	582
3	1.00	1.00	1.00	883
4	1.00	1.00	1.00	1179
accuracy			1.00	5505
macro avg	1.00	1.00	1.00	5505
weighted avg	1.00	1.00	1.00	5505



#Kmedoids when K = 3

```
from sklearn_extra.cluster import KMedoids
k=3

# Combine features into a single matrix
X = df1[['X_scale', 'Y_scale', 'cat_sc']]

# Fit KMedoids with the combined feature matrix
```

```

kmedoids = KMedoids(n_clusters=k).fit(X)

clusters = kmedoids.cluster_centers_
labels = kmedoids.labels_

sil_score = silhouette_score(X, labels)
db_score = davies_bouldin_score(X, labels)

print(f"Silhouette Score: {sil_score:.4f}")
print(f"Davies-Bouldin Index: {db_score:.4f}")

/usr/local/lib/python3.11/dist-packages/sklearn_extra/cluster/_k_medoids.py:329: UserWarning:
Cluster 1 is empty! self.labels_[self.medoid_indices_[1]] may not be
labeled with its corresponding cluster (1).

Silhouette Score: 0.2508
Davies-Bouldin Index: 1.2784

# Drop unwanted columns
df_2 = df1.drop(["Descript", "Resolution", "Address", "PdDistrict",
"DayOfWeek"], axis=1)

# Assign the cluster labels from K-Medoids
df_2['cluster'] = kmedoids.labels_

import plotly.express as px
figure1 = px.scatter_mapbox(df_2,
                           lat="Y", lon="X", # Latitude and
                           Longitude columns
                           color="cluster", # Color by cluster
                           hover_data=["Category", "X", "Y",
"cluster"], # Hover information
                           title="K-Medoids Clustering of Crime
Locations",
                           center=dict(lat=37.8, lon=-122.4), # Set
                           center for map (adjust if needed)
                           zoom=9, # Zoom level for the map
                           width=1100,
                           height=700,
                           mapbox_style="open-street-map") # Map
                           style

# Show the plot
figure1.show()

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

```

```

from xgboost import XGBClassifier
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

def train_and_evaluate_xgboost(df):
    X = df.drop(columns=['cluster', 'Dates'])
    y = df['cluster']

    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

    model = XGBClassifier(eval_metric='mlogloss', random_state=42)
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    print(f" Accuracy: {acc * 100:.2f}%\n")

    print(" Classification Report:")
    print(classification_report(y_test, y_pred))

    plt.figure(figsize=(12, 8))
    sns.heatmap(confusion_matrix(y_test, y_pred), annot=False,
fmt='d', cmap='Blues')
    plt.title("Confusion Matrix", fontsize=16)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

    return model


model = train_and_evaluate_xgboost(df_2)

```

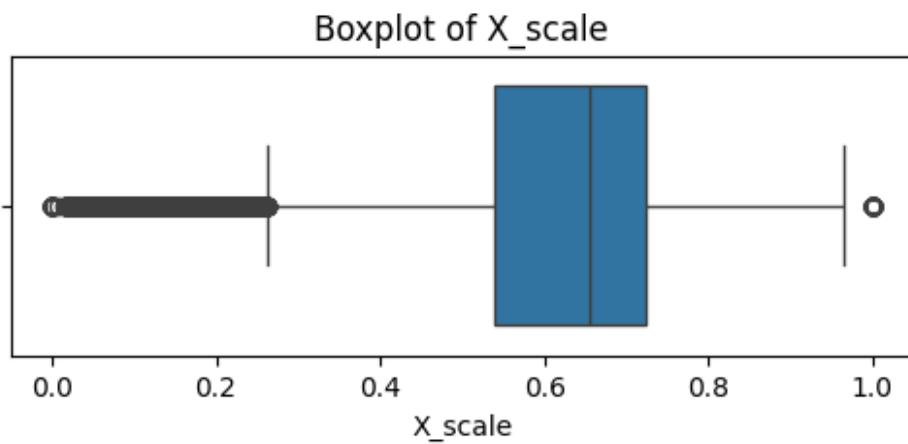
Accuracy: 99.49%

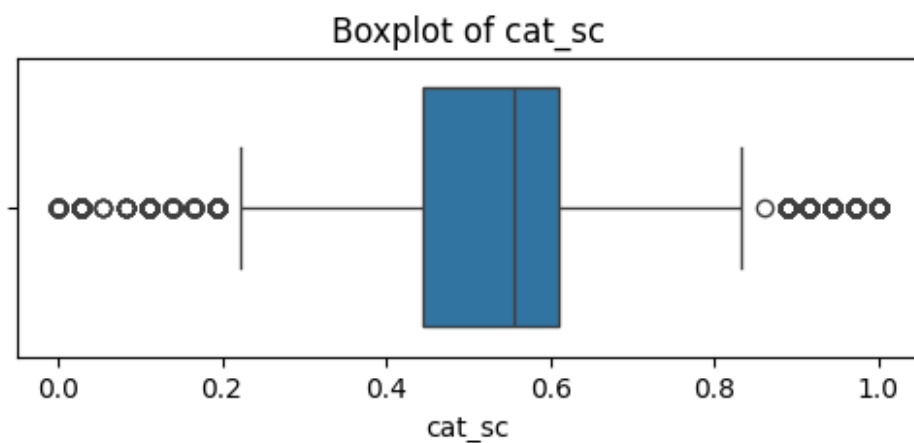

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	2430
1	0.99	0.99	0.99	1777
2	1.00	1.00	1.00	1298
accuracy			0.99	5505
macro avg	0.99	1.00	1.00	5505
weighted avg	0.99	0.99	0.99	5505



```
numerical_cols = ['X_scale', 'Y_scale', 'cat_sc']  
for col in numerical_cols:  
    plt.figure(figsize=(6, 2))  
    sns.boxplot(x=df1[col])  
    plt.title(f"Boxplot of {col}")  
    plt.show()
```



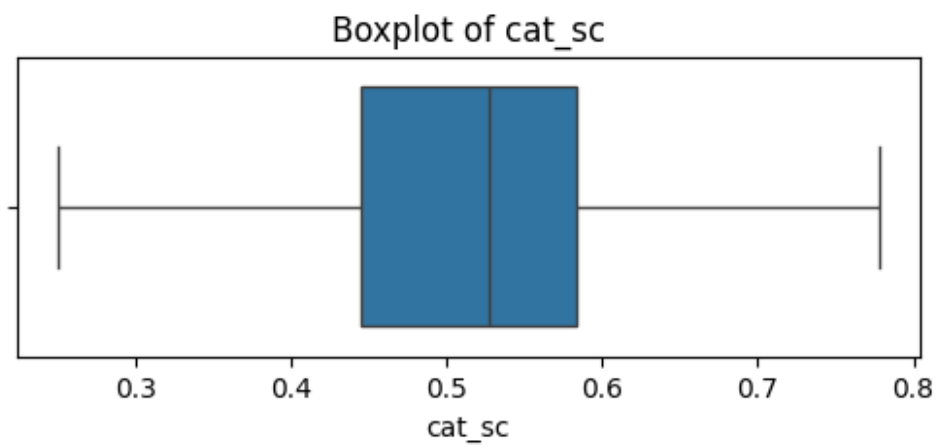
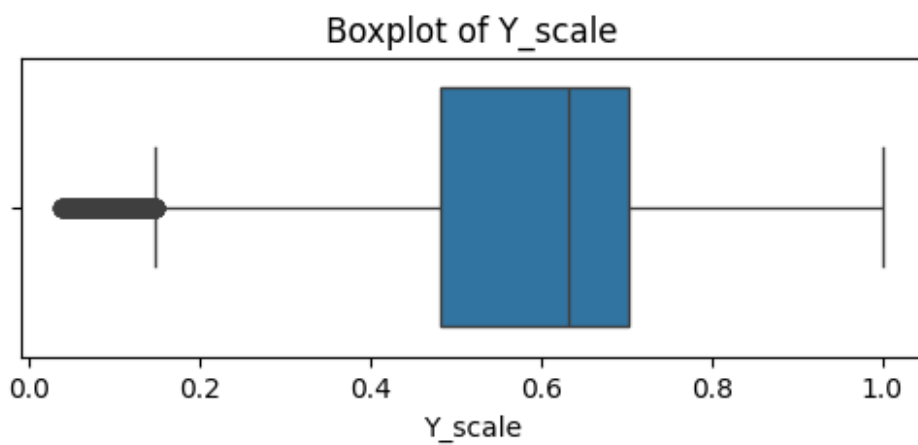
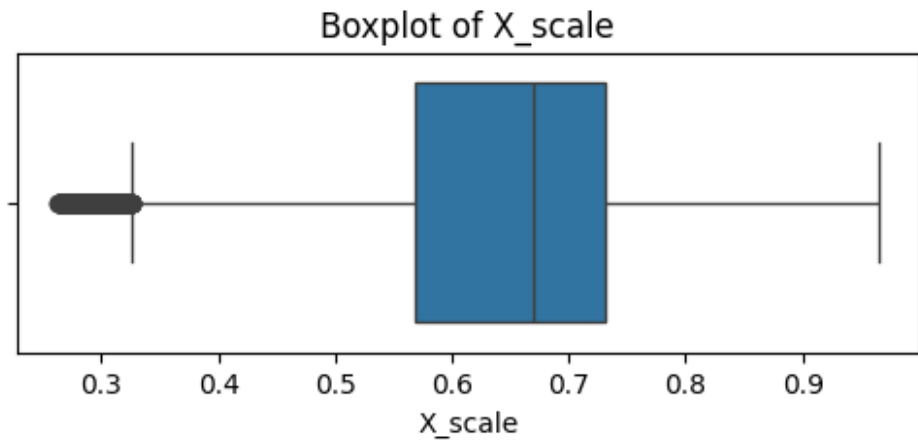


```
numerical_cols = ['X_scale', 'Y_scale', 'cat_sc']

for col in numerical_cols:
    IQR = df1[col].quantile(0.75) - df1[col].quantile(0.25)
    Lower_fence = df1[col].quantile(0.25) - (IQR * 1.5)
    Upper_fence = df1[col].quantile(0.75) + (IQR * 1.5)

    # Keep only the rows within the fences
    df1 = df1[(df1[col] >= Lower_fence) & (df1[col] <= Upper_fence)]

numerical_cols = ['X_scale', 'Y_scale', 'cat_sc']
for col in numerical_cols:
    plt.figure(figsize=(6, 2))
    sns.boxplot(x=df1[col])
    plt.title(f"Boxplot of {col}")
    plt.show()
```



#Kmedoids when K = 3 + category column

```
from sklearn_extra.cluster import KMedoids
k=3

# Combine features into a single matrix
X = df1[['X_scale', 'Y_scale', 'cat_sc']]
```

```

# Fit KMedoids with the combined feature matrix
kmedoids = KMedoids(n_clusters=k).fit(X)

clusters = kmedoids.cluster_centers_
labels = kmedoids.labels_

sil_score = silhouette_score(X, labels)
db_score = davies_bouldin_score(X, labels)

print(f"Silhouette Score: {sil_score:.4f}")
print(f"Davies-Bouldin Index: {db_score:.4f}")

/usr/local/lib/python3.11/dist-packages/sklearn_extra/cluster/_k_medoids.py:329: UserWarning:
Cluster 1 is empty! self.labels_[self.medoid_indices_[1]] may not be
labeled with its corresponding cluster (1).

Silhouette Score: 0.3216
Davies-Bouldin Index: 1.2169

# Drop unwanted columns
df_2 = df1.drop(["Descript", "Resolution", "Address", "PdDistrict",
"DayOfWeek"], axis=1)

# Assign the cluster labels from K-Medoids
df_2['cluster'] = kmedoids.labels_

import plotly.express as px
figure1 = px.scatter_mapbox(df_2,
                           lat="Y", lon="X", # Latitude and
Longitude columns
                           color="cluster", # Color by cluster
                           hover_data=["Category", "X", "Y",
"cluster"], # Hover information
                           title="K-Medoids Clustering of Crime
Locations",
                           center=dict(lat=37.8, lon=-122.4), # Set
center for map (adjust if needed)
                           zoom=9, # Zoom level for the map
                           width=1100,
                           height=700,
                           mapbox_style="open-street-map") # Map
style

# Show the plot
figure1.show()

```

```

from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix
from xgboost import XGBClassifier
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

def train_and_evaluate_xgboost(df):
    X = df.drop(columns=['Dates', 'cluster'])
    y = df['cluster']

    X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

    model = XGBClassifier(eval_metric='mlogloss', random_state=42)
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    print(f" Accuracy: {acc * 100:.2f}%\n")

    print(" Classification Report:")
    print(classification_report(y_test, y_pred))

    plt.figure(figsize=(12, 8))
    sns.heatmap(confusion_matrix(y_test, y_pred), annot=False,
fmt='d', cmap='Blues')
    plt.title("Confusion Matrix", fontsize=16)
    plt.xlabel("Predicted")
    plt.ylabel("Actual")
    plt.show()

    return model

model = train_and_evaluate_xgboost(df_2)

```

Accuracy: 99.47%

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	624
1	0.99	1.00	1.00	1674
2	0.99	0.99	0.99	906
accuracy			0.99	3204
macro avg	1.00	0.99	0.99	3204
weighted avg	0.99	0.99	0.99	3204

